



Run 5gmb8rfn: metric semantics, slow updates, and learned tool-use behavior

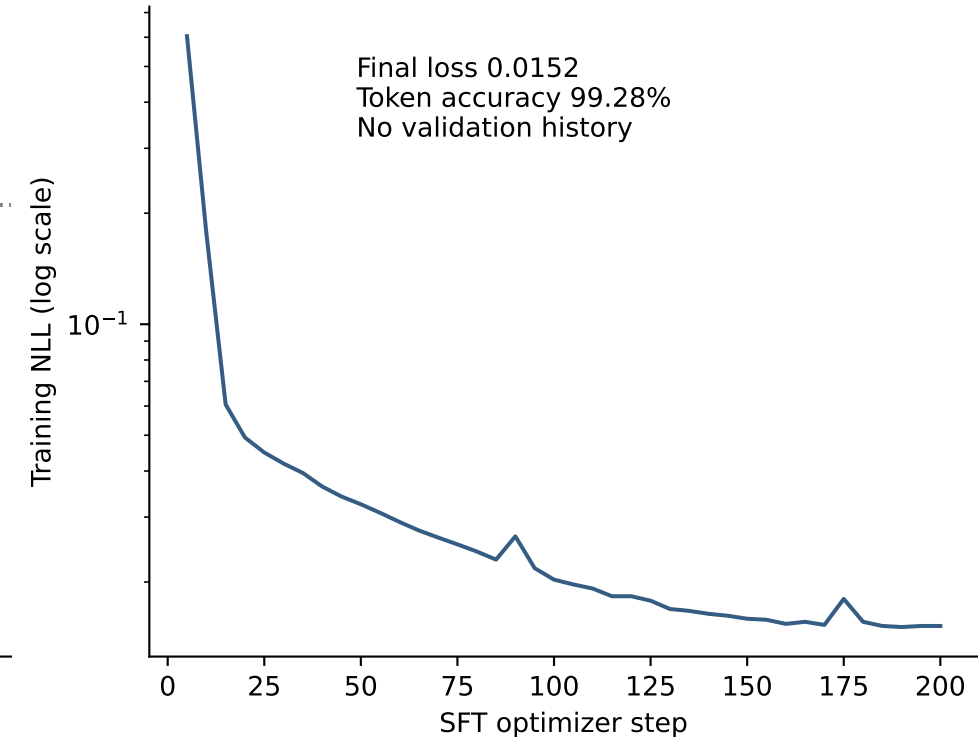
Raw detection: no demonstrated gain



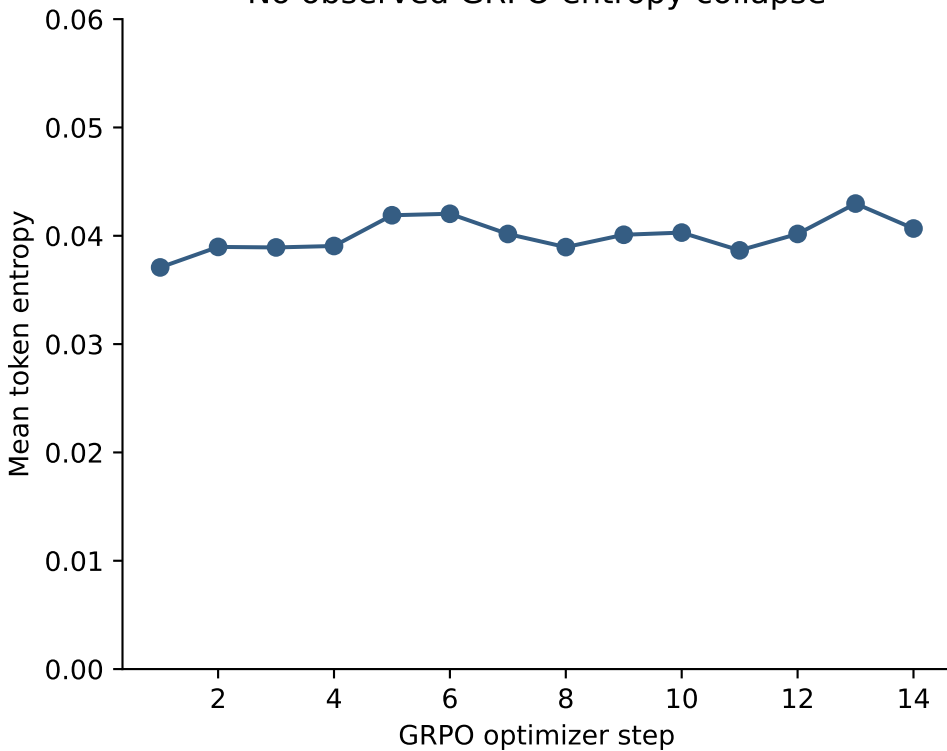
Paired change from SFT baseline



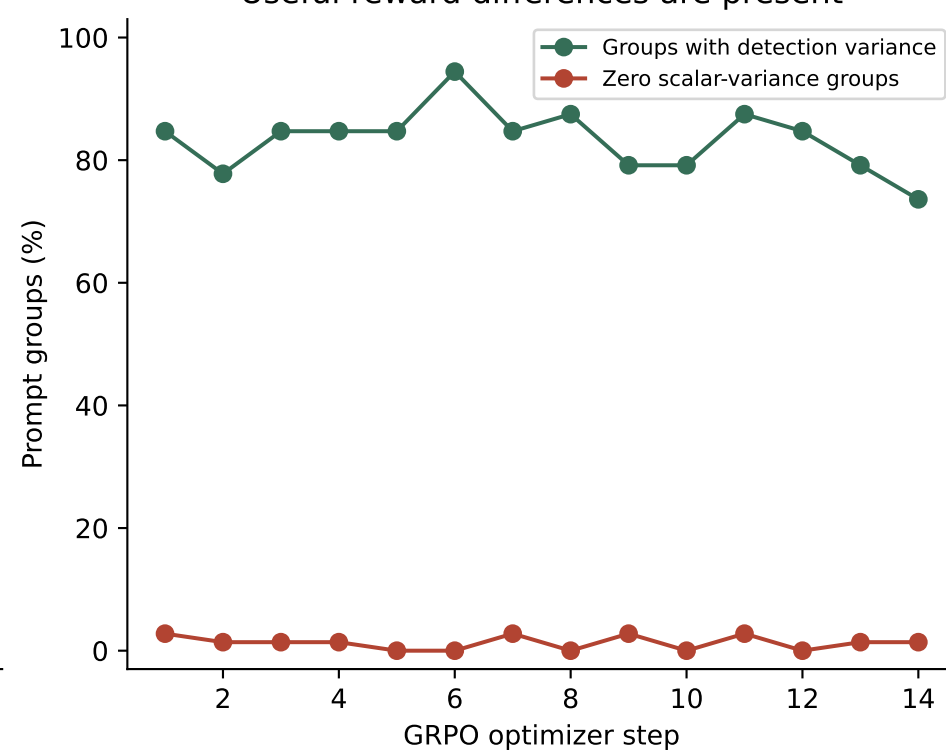
Online SFT: training loss only



No observed GRPO entropy collapse



Useful reward differences are present



Tool feedback rarely changes the answer

